

CENG 463

Introduction to Natural Language Processing

Fall 2020-2021

Assignment 2

Due date: January 17 2020, Sunday, 23:55

1 Objectives

In this assignment, you are expected to implement a series of sequence taggers. Your taggers will assign part-of-speech tags, phrase chunk labels and named-entity labels to the words in sentences. You will use 2 models, namely Logistic Regression and Conditional Random Fields for each of the 3 tasks.

You are also expected to write a report containing the analysis of your taggers and explain basic machine learning concepts. This is a research assignment and you are expected to experiment with the taggers and try different alternatives. Therefore, your reports are more important than the implementations of the taggers.

Keywords: *sequence tagging, part of speech tagging, chunking, shallow parsing, named entity recognition, maximum entropy models, logistic regression, conditional random fields*

2 Tagging Sequences

The task of labelling words in a sentence is at the beginning of almost every NLP task. Categorizing words is the most fundamental thing before moving onto more complex problems.

2.1 Part-of-Speech Tags

A part of speech of a word is the lexical category of the word (also word class, morphological class, lexical class, or lexical tag). A part of speech of a word provides information about a word and its neighbours, word's pronunciation, its morphological analysis or can be used in speech recognition, parsing, translation, information retrieval and question answering tasks. Many NLP applications may benefit from syntactically disambiguated text. Part of speech tagging is the process of assigning tags (class markers) to words. A word can be a member of multiple classes and main task of a tagger is to solve this ambiguity.

2.2 Chunks

Phrases in a sentence can be seen as blocks that move together. Finding and extracting them can help many NLP tasks. For example, a noun phrase (NP) can be an answer to a question and we can use chunking to label the words in that phrase instead of parsing the sentence to extract relations. Therefore, chunking task is also called shallow parsing. The categories include verb phrases (VP), prepositional phrases (PP), adjectives (ADJP), adverbs (ADVP) and many others.

2.3 Entities

Entity recognition tasks deal with names of persons, organizations, locations, expressions of times, quantities, monetary values etc. In this assignment, we focus on named entities which are person (PER), location (LOC) and organization (ORG) names and label some other miscellaneous names (MISC). The task again can be seen as a sequence labelling problem and solved similarly as a chunking or noun phrase chunking problem.

2.4 BIO Encoding

IOB or BIO encoding is used to assign labels in chunking or NER tasks. It is used to denote whether a word is at the beginning of a chunk/phrase/entity (B), inside it or continuation of the entity (I) or outside of any category (O). For each category, prefixes of 'B-' or 'I-' is used to label tokens. CoNLL (Conference on Natural Language Learning) datasets that are used in this assignment have BIO encoding.

3 Models

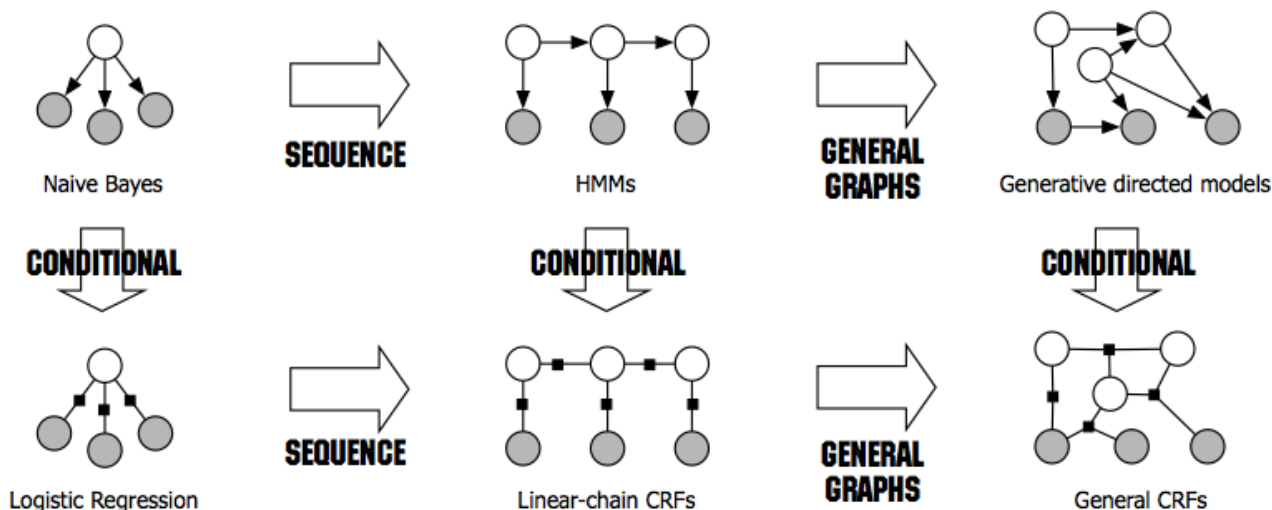


Figure 1: Diagram of relationships between models from Andrew McCallum's *An Introduction to Conditional Random Fields*

In classical machine learning, programs learn a function that maps input data to desired output labels. Data are usually represented as a set of features that are manually selected or crafted. Then a function

is learned by usually optimizing an objective function (such as minimizing error). In this assignment, we are not interested in finding the best algorithm or how to optimize the solution. We will focus on the features that are used to represent data and try to find the best set of features.

You will train two models for all 3 tasks. The first one is Maximum Entropy or Logistic Regression models and the second one is Conditional Random Fields.

3.1 Maximum Entropy Classifiers - Multinomial Logistic Regression

MaxEnt Classifiers try to predict the probability of an outcome given a set of variables. In this assignment you will represent an instance of data variable with a set of features and MaxEnt models try to find the probabilities of each class label given those features. The learning algorithm for MaxEnt tries to maximize the likelihood of the correct class for a given sample.

You will extract features from the data and train models based on these features. See the references for example features.

3.2 Conditional Random Fields

Logistic regression classifies each sample without the context it lies in. We try to add context information as additional features to the model. CRFs are more general and can take context into account to handle sequences. The sentences will be tagged as whole sequences instead of one word at a time model of MaxEnt.

You will extract very similar features in your CRF models and probably have better results than Logistic Regression models.

4 Data

There is a separate data set for each task. Each file contains tab separated fields for each token in a single line. The sentences are separated with empty lines. In each line, the first field contains the token and the last field contains the label that we want to learn and predict. Intermediate fields contains additional information that can be used to train models. These additional fields, however, will not be provided while using your models. You can train your models to predict them or use already build models to extract these features.

4.1 Part of Speech Tagging

POS data contains three fields and two types of POS tags. We are interested in predicting the POS tags in the third field. The tags in the second field are called universal part of speech tags (UPOS). UPOS is a simpler set that can be adapted to almost all languages. The tags in the third field are specific to English. Your models will predict these tags, however you can also build taggers for UPOS and use universal POS tags as features in your models.

You have `en-ud-train.conllu` and `en-ud-dev.conllu` files as training and development sections. A sample from the data can be seen below:

I	PRON	PRP
do	AUX	VBP
n't	PART	RB
think	VERB	VB
it	PRON	PRP
matters	VERB	VBZ

4.2 Chunking

The chunking data contains three fields per token. The chunks are given as B-CHUNK tag followed by zero or more I-CHUNK tags. Punctuations etc. have O tag. The second field contains POS tags derived by the Brill tagger. These are not gold (correct) POS tags. You can use these tags as features in your models, However, if you use them you also need to predict them during test time with your pos taggers (build on the POS tag data or retrained on this data) or some prebuilt taggger like Brill, Stanford, Senna or some tagger in NLTK. You can also re-tag the training data with a better tagger to improve the quality of your data. Using POS tags will affect your results dramatically.

You have only `train.txt` file to build you models. You can extract a small part as the development set or you can use cross validation. A sample from the data can be seen below:

I	NNP	B-NP
Noriega	NNP	I-NP
was	PART	B-VP
growing	VERB	I-VP
desperate	VERB	B-ADJP
.	.	O

4.3 Named Entity Recognition

The NER data contains four fields per token. The first and the last ones are the word and the entity tags. The second one is the POS tag and the third one is the chunk tag. Similarly to the chunking task you can use these as features in your models, however during the test phase only the words will be provided.

You can use `eng.train.txt` and `eng.testa.txt` files as training and development sections. A sample from the data can be seen below:

There	EX	B-NP	O
was	VBD	B-VP	O
no	DT	B-NP	O
Bundesbank	NNP	I-NP	B-ORG
intervention	NN	I-NP	O

5 Specifications

5.1 Implementation

1. You will implement six tagger classes in six modules:

Class	Module
TaggerLR_POS	<code>tagger_lr_pos.py</code>
TaggerLR_Chunk	<code>tagger_lr_chunk.py</code>
TaggerLR_NER	<code>tagger_lr_ner.py</code>
TaggerCRF_POS	<code>tagger_crf_pos.py</code>
TaggerCRF_Chunk	<code>tagger_crf_chunk.py</code>
TaggerCRF_NER	<code>tagger_crf_ner.py</code>

2. Each class will derive from the **Tagger** class in `tagger.py` module and have the methods below: *train, test, evaluate, tag, tag_sents, save, load*
See Appendix A for details.
3. You are free to add other classes, modules and methods. Given source files and class hierarchy is for guidance only and designed to minimize your work.
4. The models share almost every method with minimal modifications (mainly on feature selection).
5. Do not forget that CRF and Logistic Regression differ in their views of data. In LR, all tokens are considered separate data samples. Each word is labelled individually even though some of them are in the same sentence. In CRF, a sequence is labelled as a single entity. These will effect feature extraction and tagging.
6. Submit the best performing taggers -with parameters and feature sets- in the six tagger classes.
7. Submit additional tests, experiments and supplementary materials in other files.
8. See provided test sections in source files for usage.

5.2 Report

1. Report the following evaluation metrics.
 - precision, recall, f-score, accuracy, confusion matrix
 - Use micro averaging while reporting multi-class system scores.
 - For the NER task, precision, recall, f-score scores should omit O tags. Accuracy and confusion matrix will not.
 - Notice the high percentage of O tags, and the importance of reporting only B and I tags.
2. Explain your feature selection process
 - Rationale behind your features, why do you think they help your taggers?
 - Plot the effects of adding different features, how does the accuracy change?
 - Use 3-4 stages while adding different features and training your models. Report accuracies after adding word based, tag based, character based or position based features. One plot for each of the 3 tasks is sufficient.

3. How does data size effect your taggers?
 - Plot learning curves for a single tagger at 10% data intervals.
4. Analyze your taggers performance with the scores in the metrics above
 - Compare MaxEnt and CRF
 - Provide confusion matrices with heat map plots
 - Which classes are easy to classify, which ones are not?
5. Did you try to find better parameters for your classifiers? Report your parameter search results with plots.

6 Regulations

1. **Programming Language:** You will use Python3.
2. You can use the libraries below:
 - `nltk`: data, preprocessing and maxent
 - `scikit-learn`: maxent (logistic regression), metrics, preprocessing, feature representation etc.
 - `python-crfsuite`: CRF, feature representation and metrics
 - `sklearn-crfsuite`: scikit-learn like interface for `python-crfsuite`
 - `numpy`, `scipy`, etc.: data representation, arrays, numeric/scientific computation
 - `matplotlib`, `pandas`, `seaborn`: analysis, report and plots
3. **Late Submission** is not allowed.
4. **Cheating:** We have a zero tolerance policy for cheating. In case of a cheating event, all parts involved (source(s) and receiver(s)) get zero. People involved in cheating will be punished according to the university regulations. Remember that students of this course are bounded to the code of honour.
5. **Newsgroup:** `odtuclass`

7 Submission

- Submission will be done via `odtuclass`.
Submit a `tar.gz` file named `assgn2.tar.gz` that contains all your files related to the assignment.

8 References

- A maximum entropy model for part-of-speech tagging, Adwait Ratnaparkhi, 1996
- Conference on Natural Language Learning shared tasks:
<https://www.clips.uantwerpen.be/conll2003/ner/>

- List of the state of the art methods:
[https://aclweb.org/aclwiki/POS_Tagging_\(State_of_the_art\)](https://aclweb.org/aclwiki/POS_Tagging_(State_of_the_art))
[https://aclweb.org/aclwiki/Named_Entity_Recognition_\(State_of_the_art\)](https://aclweb.org/aclwiki/Named_Entity_Recognition_(State_of_the_art))
<https://www.clips.uantwerpen.be/conll2003/ner/>
- Documentation and download links for the libraries:
 - crfsuite
<https://sklearn-crfsuite.readthedocs.io/en/latest/>
<https://python-crfsuite.readthedocs.io/en/latest/>
 - Natural Language Toolkit
<https://www.nltk.org/>
 - scikit-learn
<https://scikit-learn.org/>

A Base Tagger Class

```

1 class Tagger():
2     '''
3     base class for all taggers
4     '''
5
6     def __init__(self, params=None):
7         '''
8         initialize the tagger (and the vectorizer if necessary)
9         params is a dict of tagger arguments that can be passed to models
10        '''
11        self.logger = None
12        self.init_logging(LOG_LEVEL)
13
14    def train(self, file_name):
15        '''
16        extract the features and train the model
17        file_name shows the path to the file with task specific format
18        '''
19        raise NotImplementedError("train not implemented")
20
21    def test(self, file_name, labels_to_remove=[]):
22        '''
23        test the model, extract features and predict tags
24        return metrics, confusion matrix and tagged data in the order below
25        precision
26        recall
27        f1
28        accuracy
29        confusion matrix
30        tagged data
31        file_name shows the path to the file with task specific format
32        labels_to_remove show classes to ignore
33        while calculating precision, recall and f1 scores
34        '''

```

```

35         raise NotImplementedError("test not implemented")
36
37     def evaluate(self, file_name, labels_to_remove=[]):
38         '''
39         return accuracy
40         to be used for validation, learning curve, parameter optimization etc.
41         '''
42         return self.test(file_name, labels_to_remove)[3]
43
44     def tag(self, sentence):
45         '''
46         tag a single tokenized sentence
47         sentence: [tkn1, tkn2, ... tknN]
48         return the tagged sentence as list of fields
49             as given in training/test files
50         tagged sentence: [[tkn1, ..., tag1], [tkn2, ..., tag2], ...]
51         put None for any field between token and tag that is not predicted
52         '''
53         raise NotImplementedError("tag not implemented")
54
55     def tag_sents(self, sentences):
56         '''
57         tag a list of tokenized sentences
58         sentences: [[tkn11, tkn12, ... tkn1N], [tkn21, tkn22, ... tkn2M] ...]
59         return the tagged sentence as a list of fields
60             as given in training/test files
61         tagged sentences: [[[tkn11, ..., tag11], ...], ...]
62         put None for any field between token and tag that is not predicted
63         '''
64         raise NotImplementedError("tag_sents not implemented")
65
66     def save(self, file_name):
67         '''
68         save the trained models for tagger to file_name
69         '''
70         raise NotImplementedError("save not implemented")
71
72     def load(self, file_name):
73         '''
74         load the trained models for tagger from file_name
75         '''
76         raise NotImplementedError("load not implemented")
77
78     def init_logging(self, log_level):
79         '''
80         logging config and init
81         '''
82         if not self.logger:
83             logging.basicConfig(
84                 format='%(asctime)s-|%(name)20s:%(funcName)12s| '
85                     '-%(levelname)8s-> %(message)s '
86                 self.logger = logging.getLogger(self.__class__.__name__)
87                 self.logger.setLevel(log_level)

```